


# Week 6 Day 1: vectorization


Notes from before:

- Issue on Quiz 5.2 (s.l?)
- Projects
- FD vs. CD for search



Memory Layout:

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$



vector of  
vectors  
(c++)



Row-major  
"C"-style  
(C, Python)



Column-major  
"F"-style  
(Fortran, matlab)

Note:

$$A^T = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$$

Also note:

$$A[\text{row}, \text{column}] \Rightarrow A = \begin{bmatrix} A_{0,0} & A_{0,1} \\ A_{1,0} & A_{1,1} \end{bmatrix}$$

but this is  $y, x$  in a graph!

C-style

$$A[a, b]$$

slow  $\nearrow$  rapid


F-style

$$A[a, b]$$


rapid  $\nearrow$  slow

# Memory Access

note: several levels



Computers love sequential access!



load to cache



asking for 0-4 is now much faster, but 5+ is still slow

Computers tend to load wide batches of memory to your cache

RAM → Cache

## Hardware vectorization

Modern hardware can compute truly in parallel on some registers!

Cache → Registers

$$\overset{A}{\begin{bmatrix} a_1 & a_2 \end{bmatrix}} + \overset{B}{\begin{bmatrix} b_1 & b_2 \end{bmatrix}} = \overset{R}{\begin{bmatrix} c_1 & c_2 \end{bmatrix}} \leftarrow \text{one instruction on CPU!}$$

You might be able to take advantage of this, but only with a proper data layout.

	32 bit	64 bit	width
SSE	4	X	128 bits
SSE2	4	2	128 bits
AVX(1+2)	8	4	256 bits
AVX512	16	8	512 bits

Take away message: if you avoid explicit loops, you'll be ready for vectorization!

Note: From now on, we'll mostly use the term "vectorization" in the MATLAB sense, for Array-at-a-time programming.

But at least now you know what real vectorization is!